

TOWARDS SAFETY & SECURITY IN C++



C-00000291*.sys

- security and safety
- safety in our code,
or lack thereof
- existing solutions for
improvement
- new features in C++26
- options for the future



ABOUT ME

- Electrical engineer
- Building computers and creating software for more than 40 years
- Developing hardware and software in the field of applied digital signal processing for more than 35 years
- Member of the C++ committee (a.k.a. WG21) for 5 years (EWG, SG15)



SOFTWARE VULNERABILITIES

- About 70 percent of Microsoft common vulnerabilities and exposures (CVEs) are memory safety vulnerabilities (based on 2006-2018 CVEs).⁸
- About 70 percent of vulnerabilities identified in Google's Chromium project are memory safety vulnerabilities.⁹
- In an analysis of Mozilla vulnerabilities, 32 of 34 critical/high bugs were memory safety vulnerabilities.¹⁰
- Based on analysis by Google's Project Zero team, 67 percent of zero-day vulnerabilities in 2021 were memory safety vulnerabilities.¹¹

The Case for Memory Safe Roadmaps, 2023

| Rank | ID | Name | Score | CVEs in KEV | Rank Change vs. 2023 |
|------|-------------------------|--|-------|-------------|----------------------|
| 1 | CWE-79 | Improper Neutralization of Input During Web Page Generation ('Cross-site Scripting') | 56.92 | 3 | +1 |
| 2 | CWE-787 | Out-of-bounds Write | 45.20 | 18 | -1 |
| 3 | CWE-89 | Improper Neutralization of Special Elements used in an SQL Command ('SQL Injection') | 35.88 | 4 | 0 |
| 4 | CWE-352 | Cross-Site Request Forgery (CSRF) | 19.57 | 0 | +5 |
| 5 | CWE-22 | Improper Limitation of a Pathname to a Restricted Directory ('Path Traversal') | 12.74 | 4 | +3 |
| 6 | CWE-125 | Out-of-bounds Read | 11.42 | 3 | +1 |
| 7 | CWE-78 | Improper Neutralization of Special Elements used in an OS Command ('OS Command Injection') | 11.30 | 5 | -2 |
| 8 | CWE-416 | Use After Free | 10.19 | 5 | -4 |
| 9 | CWE-862 | Missing Authorization | 10.11 | 0 | +2 |
| 10 | CWE-434 | Unrestricted Upload of File with Dangerous Type | 10.03 | 0 | 0 |
| 11 | CWE-94 | Improper Control of Generation of Code ('Code Injection') | 7.13 | 7 | +12 |
| 12 | CWE-20 | Improper Input Validation | 6.78 | 1 | -6 |
| 13 | CWE-77 | Improper Neutralization of Special Elements used in a Command ('Command Injection') | 6.74 | 4 | +3 |
| 14 | CWE-287 | Improper Authentication | 5.94 | 4 | -1 |
| 15 | CWE-269 | Improper Privilege Management | 5.22 | 0 | +7 |
| 16 | CWE-502 | Deserialization of Untrusted Data | 5.07 | 5 | -1 |
| 17 | CWE-200 | Exposure of Sensitive Information to an Unauthorized Actor | 5.07 | 0 | +13 |
| 18 | CWE-863 | Incorrect Authorization | 4.05 | 2 | +6 |
| 19 | CWE-918 | Server-Side Request Forgery (SSRF) | 4.05 | 2 | 0 |
| 20 | CWE-119 | Improper Restriction of Operations within the Bounds of a Memory Buffer | 3.69 | 2 | -3 |
| 21 | CWE-476 | NULL Pointer Dereference | 3.58 | 0 | -9 |
| 22 | CWE-798 | Use of Hard-coded Credentials | 3.46 | 2 | -4 |
| 23 | CWE-190 | Integer Overflow or Wraparound | 3.37 | 3 | -9 |
| 24 | CWE-400 | Uncontrolled Resource Consumption | 3.23 | 0 | +13 |
| 25 | CWE-306 | Missing Authentication for Critical Function | 2.73 | 5 | -5 |

| Rank | ID | Name | Score | CVEs in KEV | Rank Change vs. 2023 |
|------|-------------------------|--|-------|-------------|----------------------|
| 1 | CWE-79 | Improper Neutralization of Input During Web Page Generation ('Cross-site Scripting') | 56.92 | 3 | +1 |
| 2 | CWE-787 | Out-of-bounds Write | 45.20 | 18 | -1 |
| 3 | CWE-89 | Improper Neutralization of Special Elements used in an SQL Command ('SQL Injection') | 35.88 | 4 | 0 |
| 4 | CWE-352 | Cross-Site Request Forgery (CSRF) | 19.57 | 0 | +5 |
| 5 | CWE-22 | Improper Limitation of a Pathname to a Restricted Directory ('Path Traversal') | 12.74 | 4 | +3 |
| 6 | CWE-125 | Out-of-bounds Read | 11.42 | 3 | +1 |
| 7 | CWE-78 | Improper Neutralization of Special Elements used in an OS Command ('OS Command Injection') | 11.30 | 5 | -2 |
| 8 | CWE-416 | Use After Free | 10.19 | 5 | -4 |
| 9 | CWE-862 | Missing Authorization | 10.11 | 0 | +2 |
| 10 | CWE-434 | Unrestricted Upload of File with Dangerous Type | 10.03 | 0 | 0 |
| 11 | CWE-94 | Improper Control of Generation of Code ('Code Injection') | 7.13 | 7 | +12 |
| 12 | CWE-20 | Improper Input Validation | 6.78 | 1 | -6 |
| 13 | CWE-77 | Improper Neutralization of Special Elements used in a Command ('Command Injection') | 6.74 | 4 | +3 |
| 14 | CWE-287 | Improper Authentication | 5.94 | 4 | -1 |
| 15 | CWE-269 | Improper Privilege Management | 5.22 | 0 | +7 |
| 16 | CWE-502 | Deserialization of Untrusted Data | 5.07 | 5 | -1 |
| 17 | CWE-200 | Exposure of Sensitive Information to an Unauthorized Actor | 5.07 | 0 | +13 |
| 18 | CWE-863 | Incorrect Authorization | 4.05 | 2 | +6 |
| 19 | CWE-918 | Server-Side Request Forgery (SSRF) | 4.05 | 2 | 0 |
| 20 | CWE-119 | Improper Restriction of Operations within the Bounds of a Memory Buffer | 3.69 | 2 | -3 |
| 21 | CWE-476 | NULL Pointer Dereference | 3.58 | 0 | -9 |
| 22 | CWE-798 | Use of Hard-coded Credentials | 3.46 | 2 | -4 |
| 23 | CWE-190 | Integer Overflow or Wraparound | 3.37 | 3 | -9 |
| 24 | CWE-400 | Uncontrolled Resource Consumption | 3.23 | 0 | +13 |
| 25 | CWE-306 | Missing Authentication for Critical Function | 2.73 | 5 | -5 |





now

@pony.social

memory unsafety may cause a lot of software crashes, but it also helps a lot of people jailbreak their game consoles, so, it;s impossible to say if its bad or not,

Somewhere on Mastodon...

IN REAL LIFE ...





SECURITY



SECURITY



FUNCTIONAL SAFETY



FUNCTIONAL SAFETY



GOING DOWN ...


```
1 // contents of f.cpp
2
3 #include "f.h"
4
5 constexpr int f() {
6     S s;
7     U u;
8
9     int *sp = s.b;
10    int *sq = &s.b[0];
11    int *sr = &s.b[4];
12    int *st = &s.c;
13
14    assert( sp == sq);
15    assert( sr == st);
16    assert(*sr == *st);
17
18    int *up = u.b;
19    int *uq = &u.b[0];
20    int *ut = &u.c;
21
22    assert( up == uq);
23    assert( uq == ut);
24    assert(*uq == *ut);
25
26    return 0;
27 }
28
29 int main() {
30     int rc = f();
31     return rc;
32 }
```

```
1 // contents of f.h
2
3 // product type
4
5 struct S {
6     int a;
7     int b[4];
8     int c;
9 };
10
11 // sum type
12
13 union U {
14     int a;
15     int b[4];
16     int c;
17 };
```

<https://godbolt.org/z/Tc13f7YbT>


```
1 // contents of f.cpp
2
3 #include "f.h"
4
5 constexpr int f() {
6     S s{};           // no EB or UB =>
7     U u{};           // 👍 at compile-time
8
9     int *sp = s.b;
10    int *sq = &s.b[0];
11    int *sr = &s.b[4];
12    int *st = &s.c;
13
14    assert( sp == sq);
15    assert( sr == st);
16    assert(*sr == *st);
17
18    int *up = u.b;
19    int *uq = &u.b[0];
20    int *ut = &u.c;
21
22    assert( up == uq);
23    assert( uq == ut);
24    assert(*uq == *ut);
25
26    return 0;
27 }
28
29 int main() {
30     constexpr int rc = f();
31     return rc;
32 }
```

```
1 // contents of f.h
2
3 // product type
4
5 struct S {
6     int a;
7     int b[4];
8     int c;
9 };
10
11 // sum type
12
13 union U {
14     int a;
15     int b[4];
16     int c;
17 };
```

<https://godbolt.org/z/Trzo7rcPG>

```
1 // contents of f.cpp
2
3 #include "f.h"
4
5 constexpr int f() {
6     S s{};           // no EB or UB ⇒
7     U u{};           // 👍 at compile-time
8
9     int *sp = s.b;
10    int *sq = &s.b[0];
11    int *sr = &s.b[4];
12    int *st = &s.c;
13
14    assert( sp == sq);
15    assert( sr == st);
16    assert(*sr == *st);
17
18    int *up = u.b;
19    int *uq = &u.b[0];
20    int *ut = &u.c;
21
22    assert( up == uq);
23    assert( uq == ut);
24    assert(*uq == *ut);
25
26    return 0;
27 }
28
29 int main() {
30     constexpr int rc = f();
31     return rc;
32 }
```

```
1 // contents of f.h
2
3 // product type
4
5 struct S {
6     int a;
7     int b[4];
8     int c;
9 };
10
11 // sum type
12
13 union U {
14     int a;
15     int b[4];
16     int c;
17 };
```

<https://godbolt.org/z/Trzo7rcPG>

The current C++ standard document
mentions "undefined behaviour" (UB)
≈ 90 times in the language-specific part.





Constant evaluation

Constant evaluation
is mandated to
refuse and diagnose all occurrences of UB
during program execution.



Constant evaluation
is mandated to
refuse and diagnose all occurrences of UB
during program execution.

[defns.undefined], [expr.const]/10

CONSTEXPR

- since C++11
- major improvements until C++20
- minor language improvements in C++23,
but huge library improvements

CONSTEXPR

- since C++11
- major improvements until C++20
- minor language improvements in C++23,
but huge library improvements

new in C++26:

- exceptions
- placement new
- structured bindings
- improved references
- lots more library functions





MORE CODE ...

```
1 constexpr int f(const int *array, bool a, bool b) {
2     unsigned index = 1;
3
4     if (a) --index;
5     if (b) --index;
6
7     return array[index];
8 }
9
10 constexpr int array[4]{};
11
12 static_assert(f(array, false, false) == 0);
```



```
1 constexpr int f(const int *array, bool a, bool b) {
2     unsigned index = 1;
3
4     if (a) --index;
5     if (b) --index;
6
7     // implicit contract: index lies within the referenced array
8     return array[index];
9 }
10
11 constexpr int array[4]{};
12
13 static_assert(f(array, false, false) == 0);
```



```
1 constexpr int f(const int *array, bool a, bool b) {
2     unsigned index = 1;
3
4     if (a) --index;
5     if (b) --index;
6
7     // implicit contract: index lies within the referenced array
8     return array[index];
9 }
10
11 constexpr int array[4] {};
12
13 static_assert(f(array, false, false) == 0);
```



```
1 constexpr int f(const int *array, bool a, bool b) {
2     unsigned index = 1;
3
4     if (a) --index;
5     if (b) --index;
6
7     // implicit contract: index lies within the referenced array
8     return array[index];
9 }
10
11 constexpr int array[4]{};
12
13 static_assert(f(array, false, false) == 0);
```



```
1 constexpr int f(const int *array, bool a, bool b) {
2     unsigned index = 1;
3
4     if (a) --index;
5     if (b) --index;
6
7     // implicit contract: index lies within the referenced array
8     return array[index];
9 }
10
11 constexpr int array[4]{};
12
13 static_assert(f(array, false, false) == 0);
```



```
1 constexpr int f(const int *array, bool a, bool b) {
2     unsigned index = 1;
3
4     if (a) --index;
5     if (b) --index;
6
7     // implicit contract: index lies within the referenced array
8     return array[index];
9 }
10
11 constexpr int array[4]{};
12
13 static_assert(f(array, false, false) == 0);
```

```
1 constexpr int f(const int (&array)[4], bool a, bool b) {
2     unsigned index = 1;
3
4     if (a) --index;
5     if (b) --index;
6
7     // perform a range check here!
8     return array[index];
9 }
10
11 constexpr int array[4]{};
12
13 static_assert(f(array, false, false) == 0);
```



```
1 template <std::size_t N>
2 constexpr int f(const int (&array)[N], bool a, bool b) {
3     unsigned index = 1;
4
5     if (a) --index;
6     if (b) --index;
7
8     // perform a range check here!!
9     return array[index];
10 }
11
12 constexpr int array[4]{};
13
14 static_assert(f(array, false, false) == 0);
```



```
1 constexpr int f(const int array[], std::size_t N, bool a, bool b) {
2     unsigned index = 1;
3
4     if (a) --index;
5     if (b) --index;
6
7     // perform a range check here!!!!
8     return array[index];
9 }
10
11 constexpr int array[4]{};
12
13 static_assert(f(array, std::size(array), false, false) == 0);
```



```
1 #include <span> // "Contemporary C++" -> use strong types
2
3 constexpr int select(std::span<const int> sequence,
4                      SomeBooleanOption a, AnotherBooleanOption b) {
5     unsigned index = 1;
6
7     if (a) --index;
8     if (b) --index;
9
10    // !!! perform a range check here !!!
11    return sequence[index];
12 }
13
14 constexpr int array[4]{};
15
16 static_assert(select(array, option_a, option_b) == 0);
```

```
1 #include <span> // "Contemporary C++" -> use strong types
2
3 constexpr int select(std::span<const int> sequence,
4                      SomeBooleanOption a, AnotherBooleanOption b) {
5     unsigned index = 1;
6
7     if (a) --index;
8     if (b) --index;
9
10    // !!! perform a range check here !!!
11    return sequence[index];
12 }
13
14 constexpr int array[4]{};
15
16 static_assert(select(array, option_a, option_b) == 0);
```



```
1 #include <span>
2
3 constexpr int select(std::span<const int> sequence,
4                      SomeBooleanOption a, AnotherBooleanOption b) {
5     unsigned index = 1;
6
7     if (a) --index;
8     if (b) --index;
9
10    // range checking is internally done
11    // in a 'hardened' standard library!
12    return sequence[index];
13 }
14
15 constexpr int array[4]{};
16
17 static_assert(select(array, option_a, option_b) == 0);
```



P3471

Standard library hardening

P3471



Standard library hardening



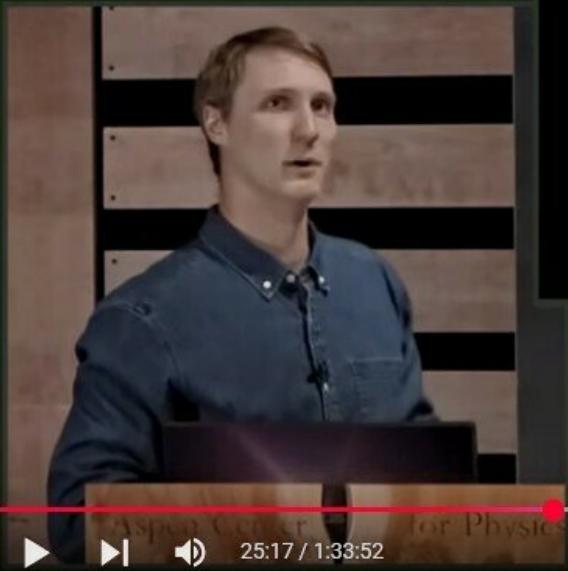
Expect this even before C++26 !



CppNow.org

Video Sponsorship Provided By

millennium
think-cell



Agenda

Overview of Memory Safety

Library Undefined Behavior

Standard Library Hardening

Typed Memory Operations

Conclusions

i

27

Security in C++
Hardening Techniques From the Trenches

Louis Dionne



Louis Dionne "Security in C++", C++Now 2024

C++ library preconditions

- are either checked with termination semantics in all compilation modes,
- or are ignored → UB (status quo)
- detected contract violations
 - may stop execution instantly (without P2900)
 - detour to handle_contractViolation (with P2900)

C++ library preconditions

- are either checked with termination semantics in all compilation modes,
 - or are ignored → UB (status quo)
-
- detected contract violations
 - may stop execution instantly (without P2900)
 - detour to handle_contractViolation (with P2900)



C++ standard library implementations may go **beyond** P3471

- more containers
- more classes
- more functions
- destructor pointer tombstones → resist "use after free" 😎 🛡️

C++ standard library implementations may go **beyond** P3471

- more containers
- more classes
- more functions
- destructor pointer tombstones → resist "use after free" 😎 🛡️

Available in MS-STL, libc++, ... !

```
1 import std;
2
3 int main(int argc, char*[]) {
4     std::println("hardened STL: {}", static_cast<bool>(_MSVC_STL_HARDENING));
5
6     std::vector v(std::max(argc, 1), argc);
7     std::println("v[3] = {}", v[3]);
8     std::println("v[-1] = {}", v[-1]);
9 }
```

MS-STL version 202502

just as ever

with hardening

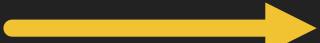
```
1 import std;
2
3 int main(int argc, char*[]) {
4     std::println("hardened STL: {}", static_cast<bool>(_MSVC_STL_HARDENING));
5
6     std::vector v(std::max(argc, 1), argc);
7     std::println("v[3] = {}", v[3]);
8     std::println("v[-1] = {}", v[-1]);
9 }
```

MS-STL version 202502

just as ever

with hardening

```
1 > cl test.cpp
2 > test.exe
3 hardened STL: false
4 v[3] = 9673
5 v[-1] = -1811929271
6
7 // program ended with return value 0
```



```
1 import std;
2
3 int main(int argc, char*[]) {
4     std::println("hardened STL: {}", static_cast<bool>(_MSVC_STL_HARDENING));
5
6     std::vector v(std::max(argc, 1), argc);
7     std::println("v[3] = {}", v[3]);
8     std::println("v[-1] = {}", v[-1]);
9 }
```

MS-STL version 202502

just as ever

```
1 > cl test.cpp
2 > test.exe
3 hardened STL: false
4 v[3] = 9673
5 v[-1] = -1811929271
6
7 // program ended with return value 0
```



with hardening

```
1 > cl test.cpp
2 > test.exe
3 hardened STL: true
4
5 // program execution stopped here
6 // with error code 0xC0000409
7 // i.e. STATUS_FAIL_FAST_EXCEPTION
```



DILIGENCE



P2900

Contracts for C++



P2900

Contracts for C++

Similar to Eiffel, Ada, or D

CONTRACTS IN C++

Support Developers

CONTRACTS IN C++

Assist "Design by Contract"

CONTRACTS IN C++

Check Correctness

```
1 // somewhere in the program source
2
3 int func(int x)
4 // precondition: x is not 1
5 // postcondition: result value is different from argument value
6 {
7 // ...
8     assert(x != 3); // assures an expectation
9 // ...
10    return x | 2;
11 }
12
13 void call_func() {
14     func(0); // fine
15     func(1); // oops
16     func(2); // oh my...
17     func(3); // what now?
18     func(4); // fine again
19 }
```



```
1 // somewhere in the program source
2
3 int func(int x)
4 // precondition: x is not 1
5 // postcondition: result value is different from argument value
6 {
7 // ...
8     assert(x != 3); // assures an expectation
9 // ...
10    return x | 2;
11 }
12
13 void call_func() {
14     func(0); // fine
15     func(1); // oops
16     func(2); // oh my...
17     func(3); // what now?
18     func(4); // fine again
19 }
```

```

1 // somewhere in the program source
2
3 int func(int x)           // function parameters are *immutable* in pre/post assertions!
4     pre (x != 1)          // a precondition assertion
5     post(r : r != x)      // a postcondition assertion
6                     // 'r' names the result object of 'func'
7 {
8     // observable checkpoint (P1494 ⚡) right before evaluating the assert-argument
9     contract_assert(x != 3); // an assertion statement, e.g. check invariants
10    // observable checkpoint right after returning from the handler in 'observe' semantics
11    return x | 2;
12 }
13
14 void call_func() {
15     func(0);   // no contract violation
16     func(1);   // violates precondition assertion of 'func'
17     func(2);   // violates postcondition assertion of 'func'
18     func(3);   // violates assertion statement within 'func'
19     func(4);   // no contract violation
20 }
21
22 // -----
23
24 // in global namespace, global module, invisible to users
25
26 // called on any contract violation
27 // system-provided or optionally user-replaceable
28
29 void handle_contractViolation(std::contracts::contractViolation cvo) {
30     // handle contract violation according to
31     // - the evaluation semantics as indicated in 'cvo'
32     // - the rest of the info in 'cvo'
33 }
```



```
1 // somewhere in the program source
2
3 int func(int x)          // function parameters are *immutable* in pre/post assertions!
4     pre (x != 1)          // a precondition assertion
5     post(r : r != x)      // a postcondition assertion
6                     // 'r' names the result object of 'func'
7 {
8     // observable checkpoint (P1494 ⚡) right before evaluating the assert-argument
9     contract_assert(x != 3); // an assertion statement, e.g. check invariants
10    // observable checkpoint right after returning from the handler in 'observe' semantics
11    return x | 2;
12 }
13
14 void call_func() {
15     func(0);   // no contract violation
16     func(1);   // violates precondition assertion of 'func'
17     func(2);   // violates postcondition assertion of 'func'
18     func(3);   // violates assertion statement within 'func'
19     func(4);   // no contract violation
20 }
21
22 // -----
23
24 // in global namespace, global module, invisible to users
25
26 // called on any contract violation
27 // system-provided or optionally user-replaceable
28
29 void handle_contractViolation(std::contracts::contractViolation cvo) {
30     // handle contract violation according to
31     // - the evaluation semantics as indicated in 'cvo'
32     // - the rest of the info in 'cvo'
33 }
```

CONTRACT EVALUATION SEMANTICS

| Semantics | check condition | fail → invoke handler | terminate |
|------------------------|-----------------|-----------------------|----------------|
| ignore (zero overhead) | ✗ | ✗ | ✗ |
| observe | ✓ | ✓ ¹ | ✗ |
| enforce | ✓ | ✓ | ✓ ² |
| quick_enforce | ✓ | ✗ | ✓ ³ |

¹ runtime: proceed after calling the handler / compile-time: diagnostic

² runtime: std::abort / compile-time: program is ill-formed

³ runtime: terminate abstract machine execution (i.e. no more user code)
compile-time: program is ill-formed

CONTRACT EVALUATION SEMANTICS

| Semantics | check condition | fail → invoke handler | terminate |
|------------------------|-----------------|-----------------------|----------------|
| ignore (zero overhead) | ✗ | ✗ | ✗ |
| observe | ✓ | ✓ ¹ | ✗ |
| enforce | ✓ | ✓ | ✓ ² |
| quick_enforce | ✓ | ✗ | ✓ ³ |

¹ runtime: proceed after calling the handler / compile-time: diagnostic

² runtime: std::abort / compile-time: program is ill-formed

³ runtime: terminate abstract machine execution (i.e. no more user code)
compile-time: program is ill-formed





SAFE C++ ?

P3081

Core safety profiles for C++26

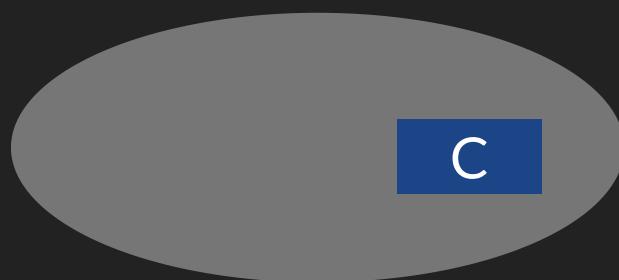


P3081

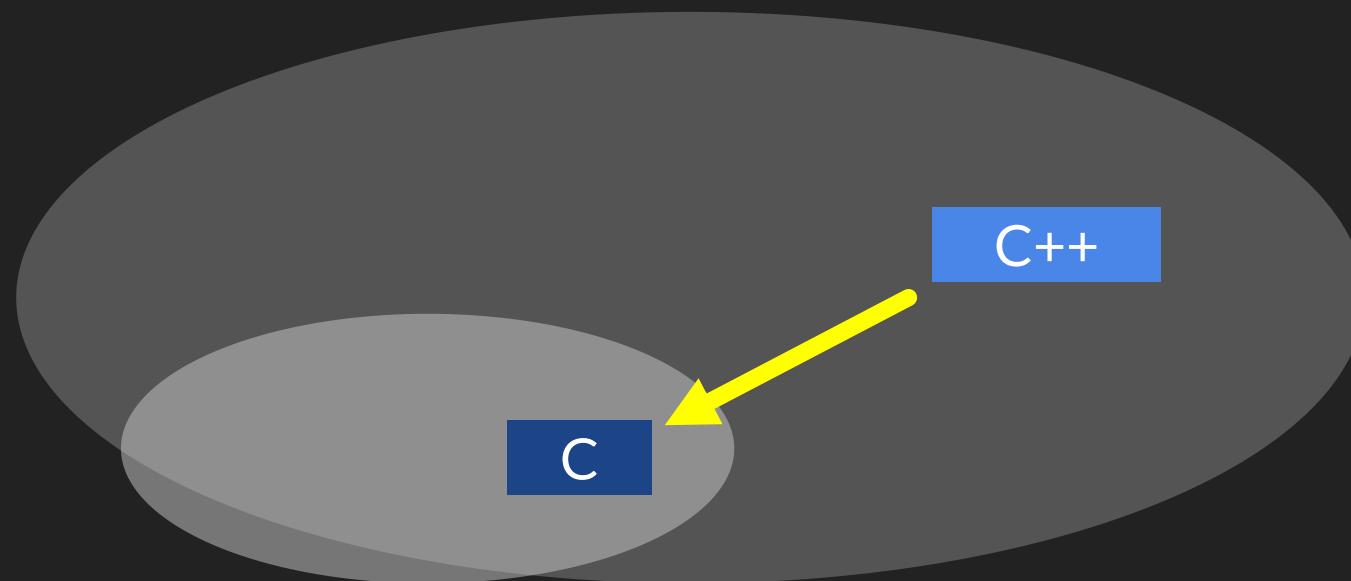
Core safety profiles for C++26

Not in C++26 ! 😬

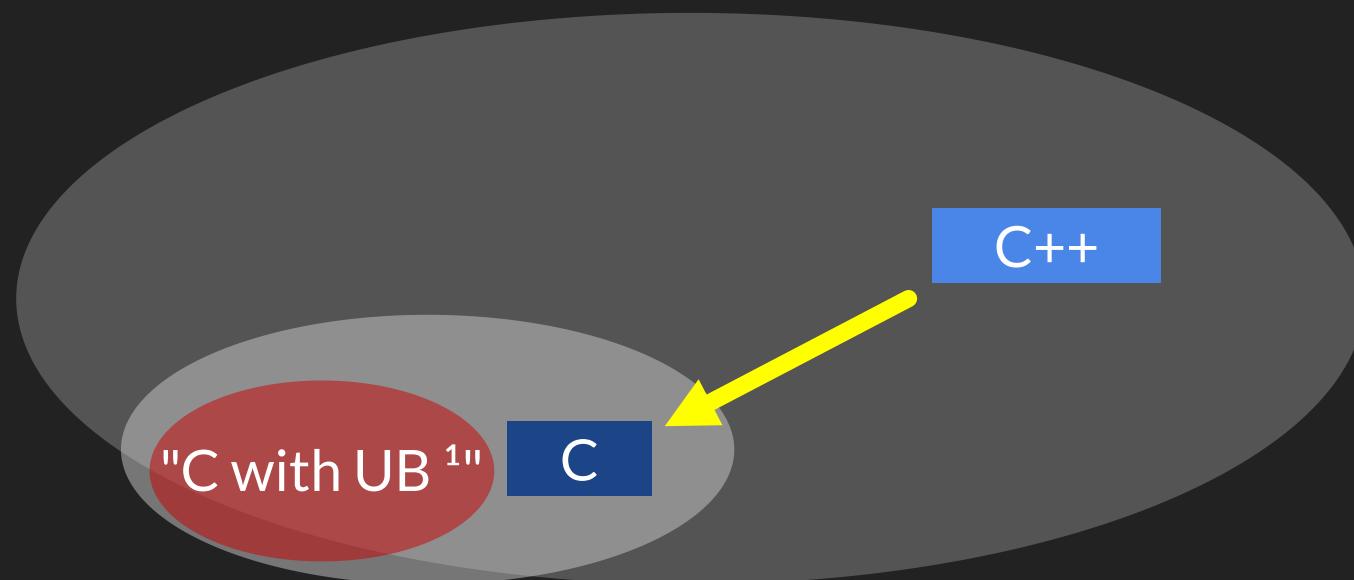
THE IDEA



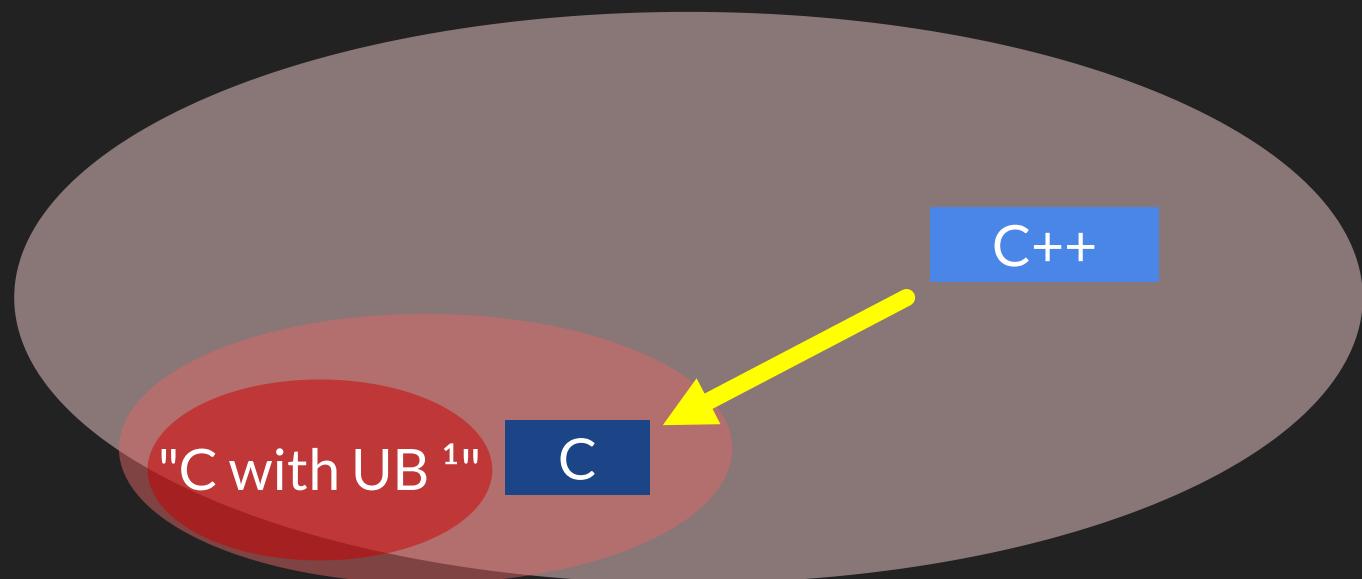
THE IDEA



THE IDEA



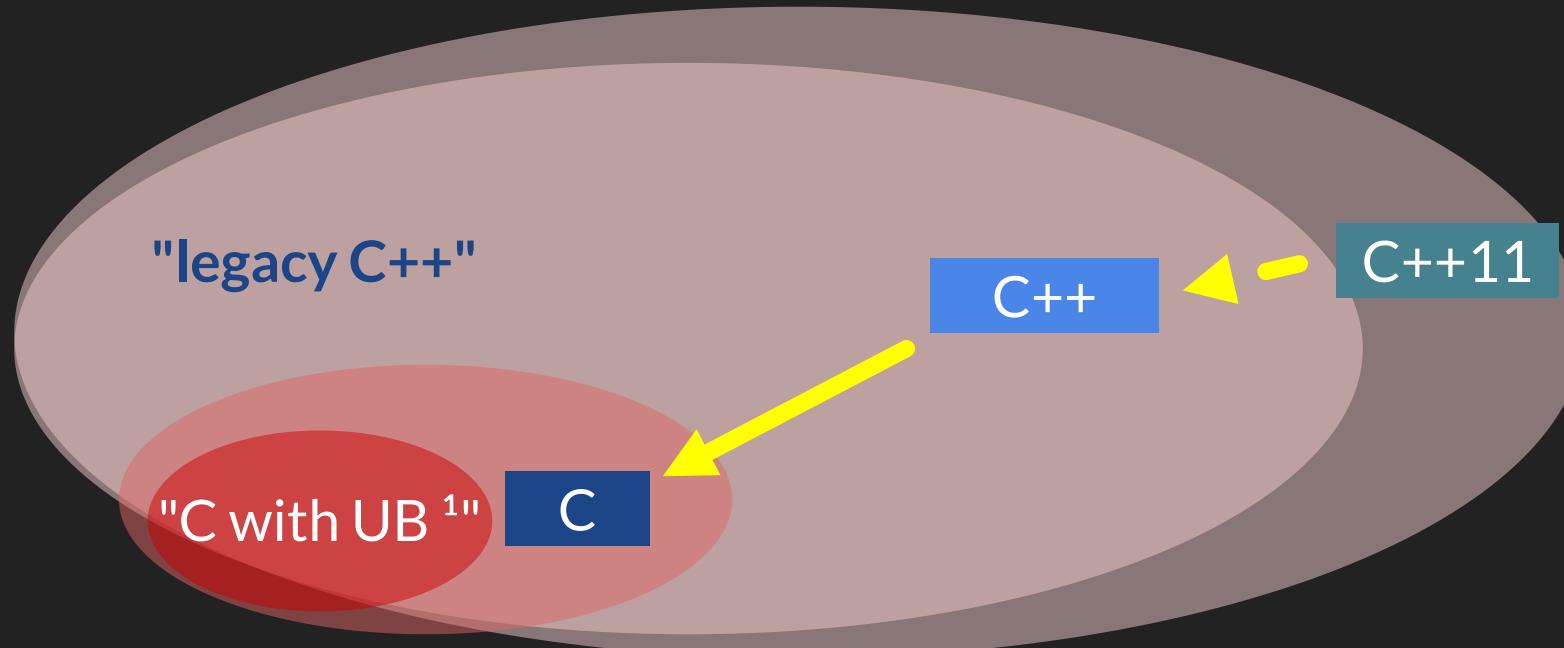
THE IDEA



¹ WG14 N3453: C contains

- 100 cases of UB in the core language
- 221 cases of UB in total

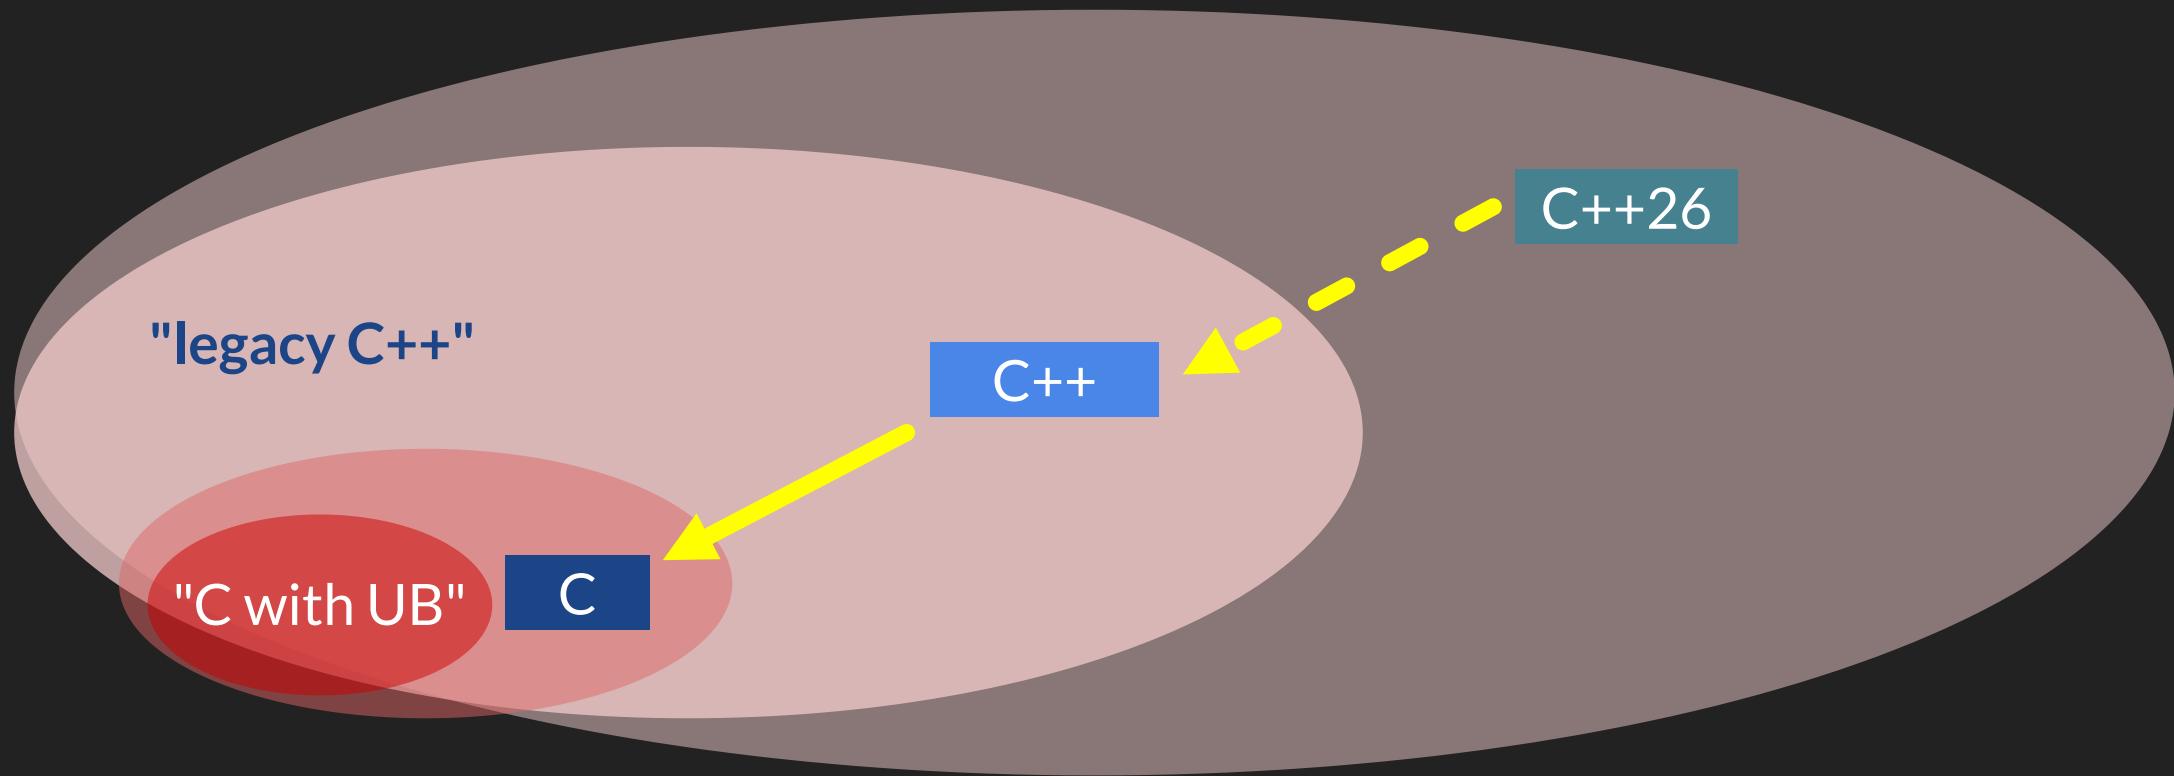
THE IDEA



¹ WG14 N3453: C contains

- 100 cases of UB in the core language
- 221 cases of UB in total

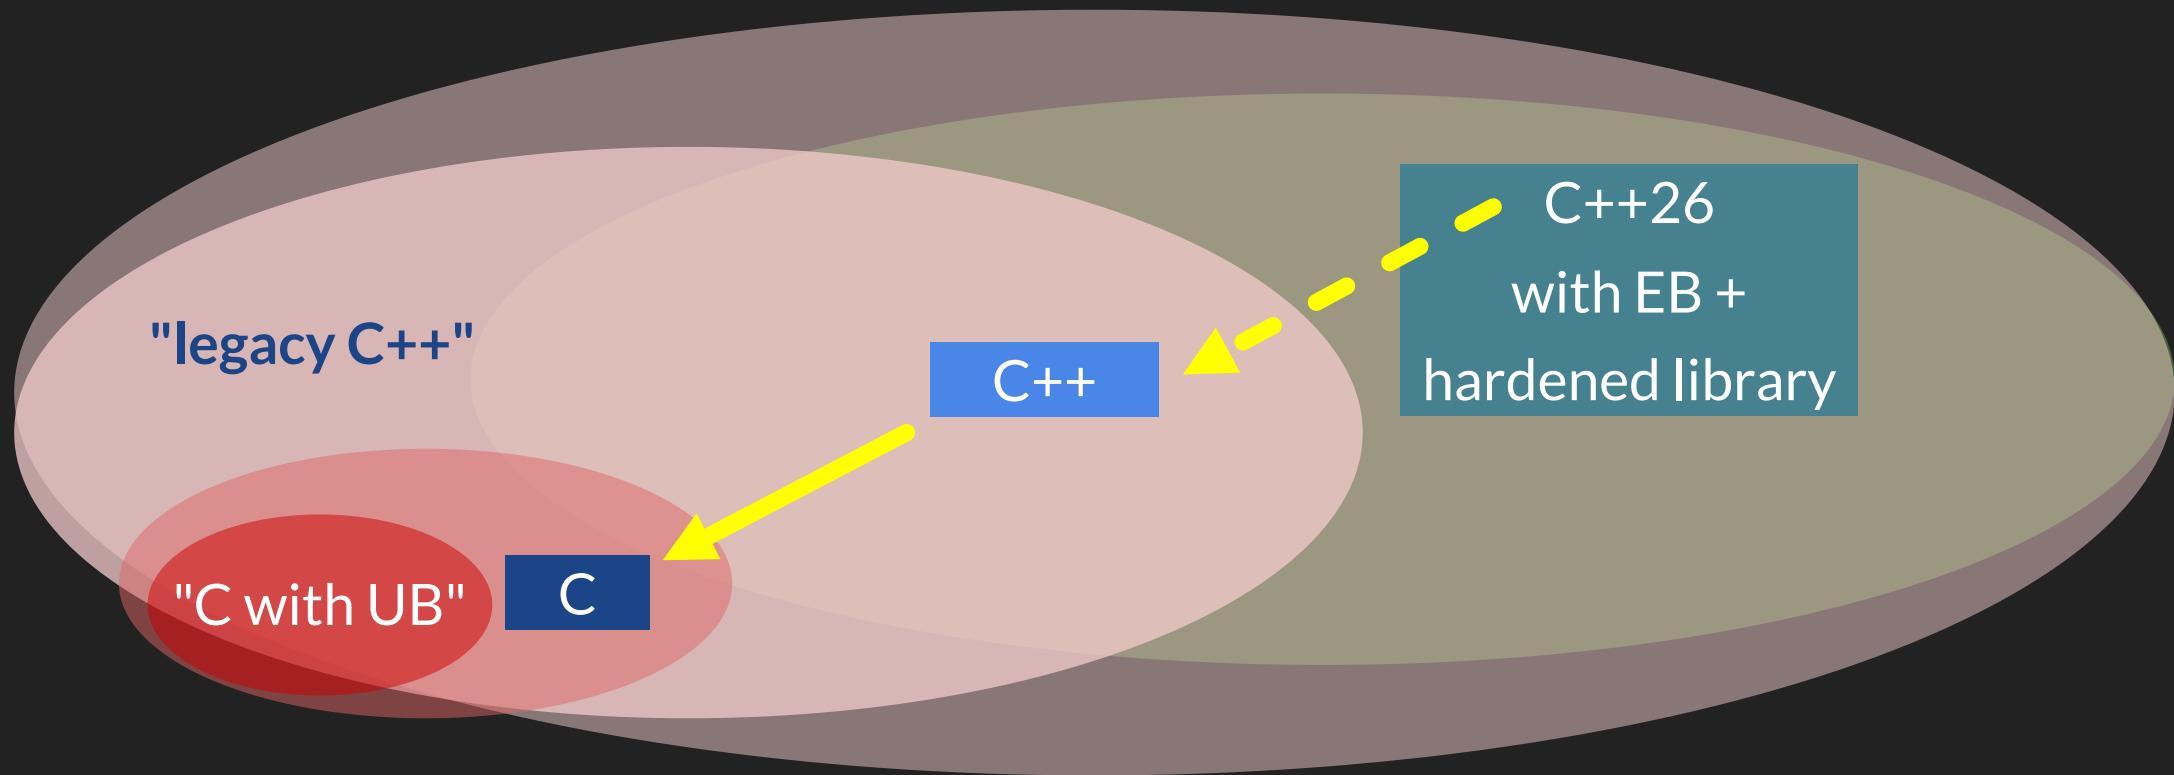
THE IDEA



Inherited UB taints the entire language

C++ adds its own UB

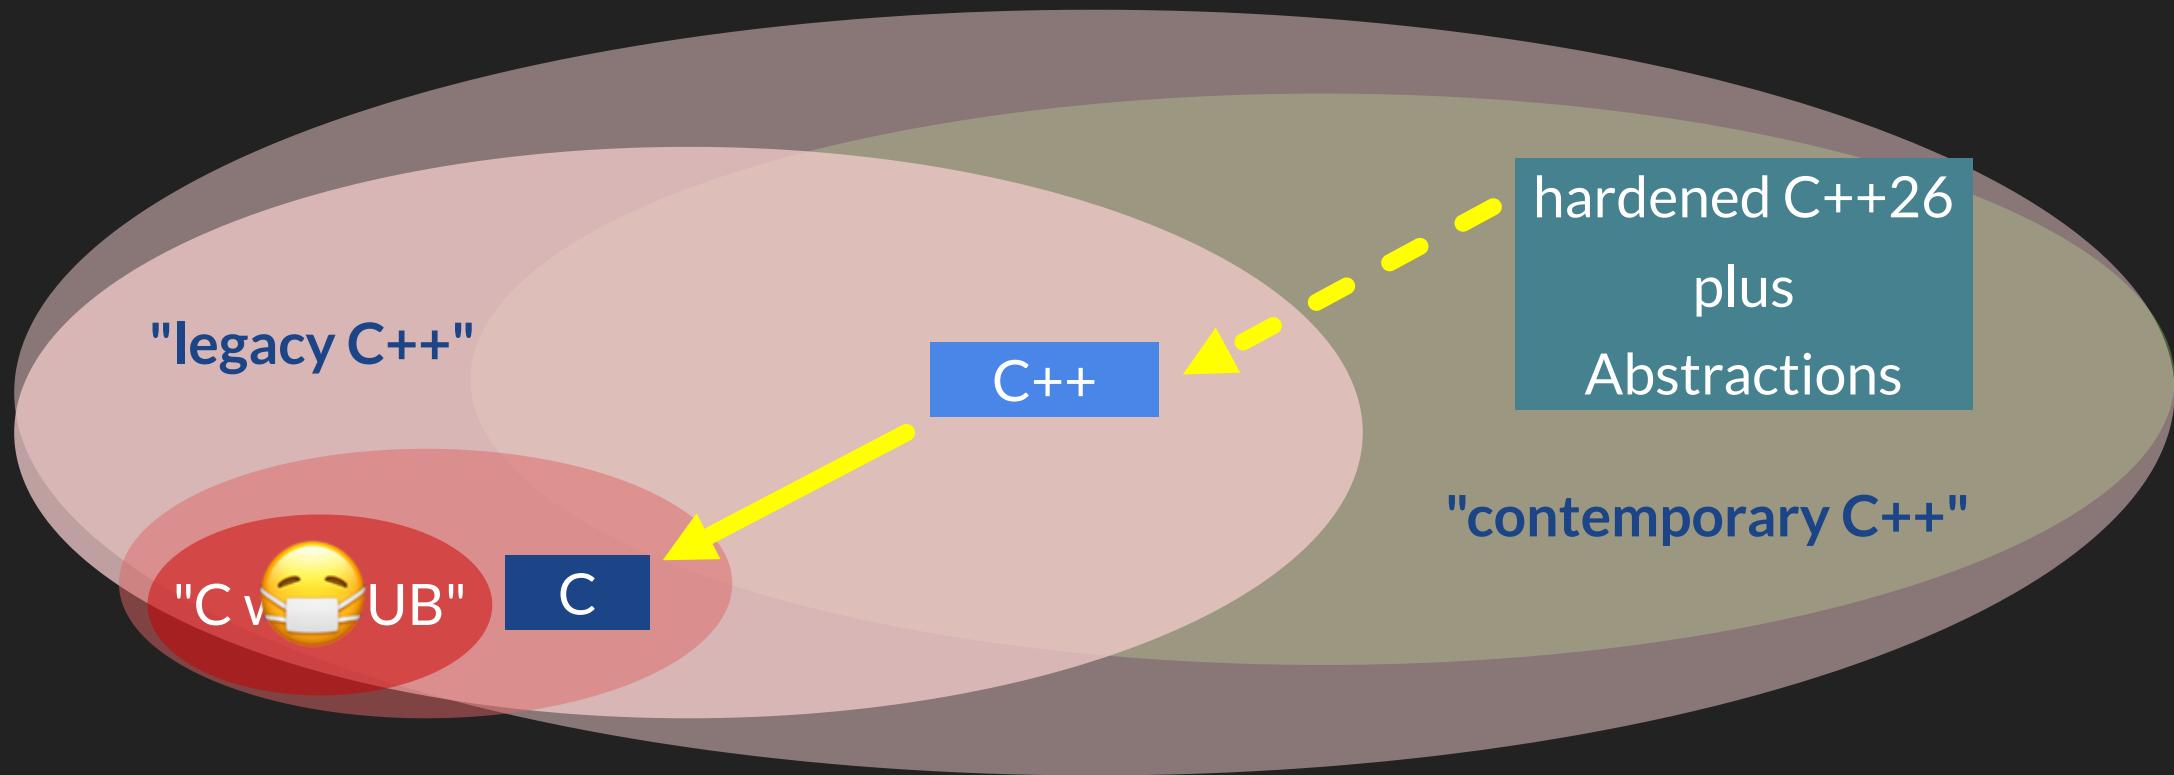
THE IDEA



Inherited UB taints the entire language

C++ adds its own UB

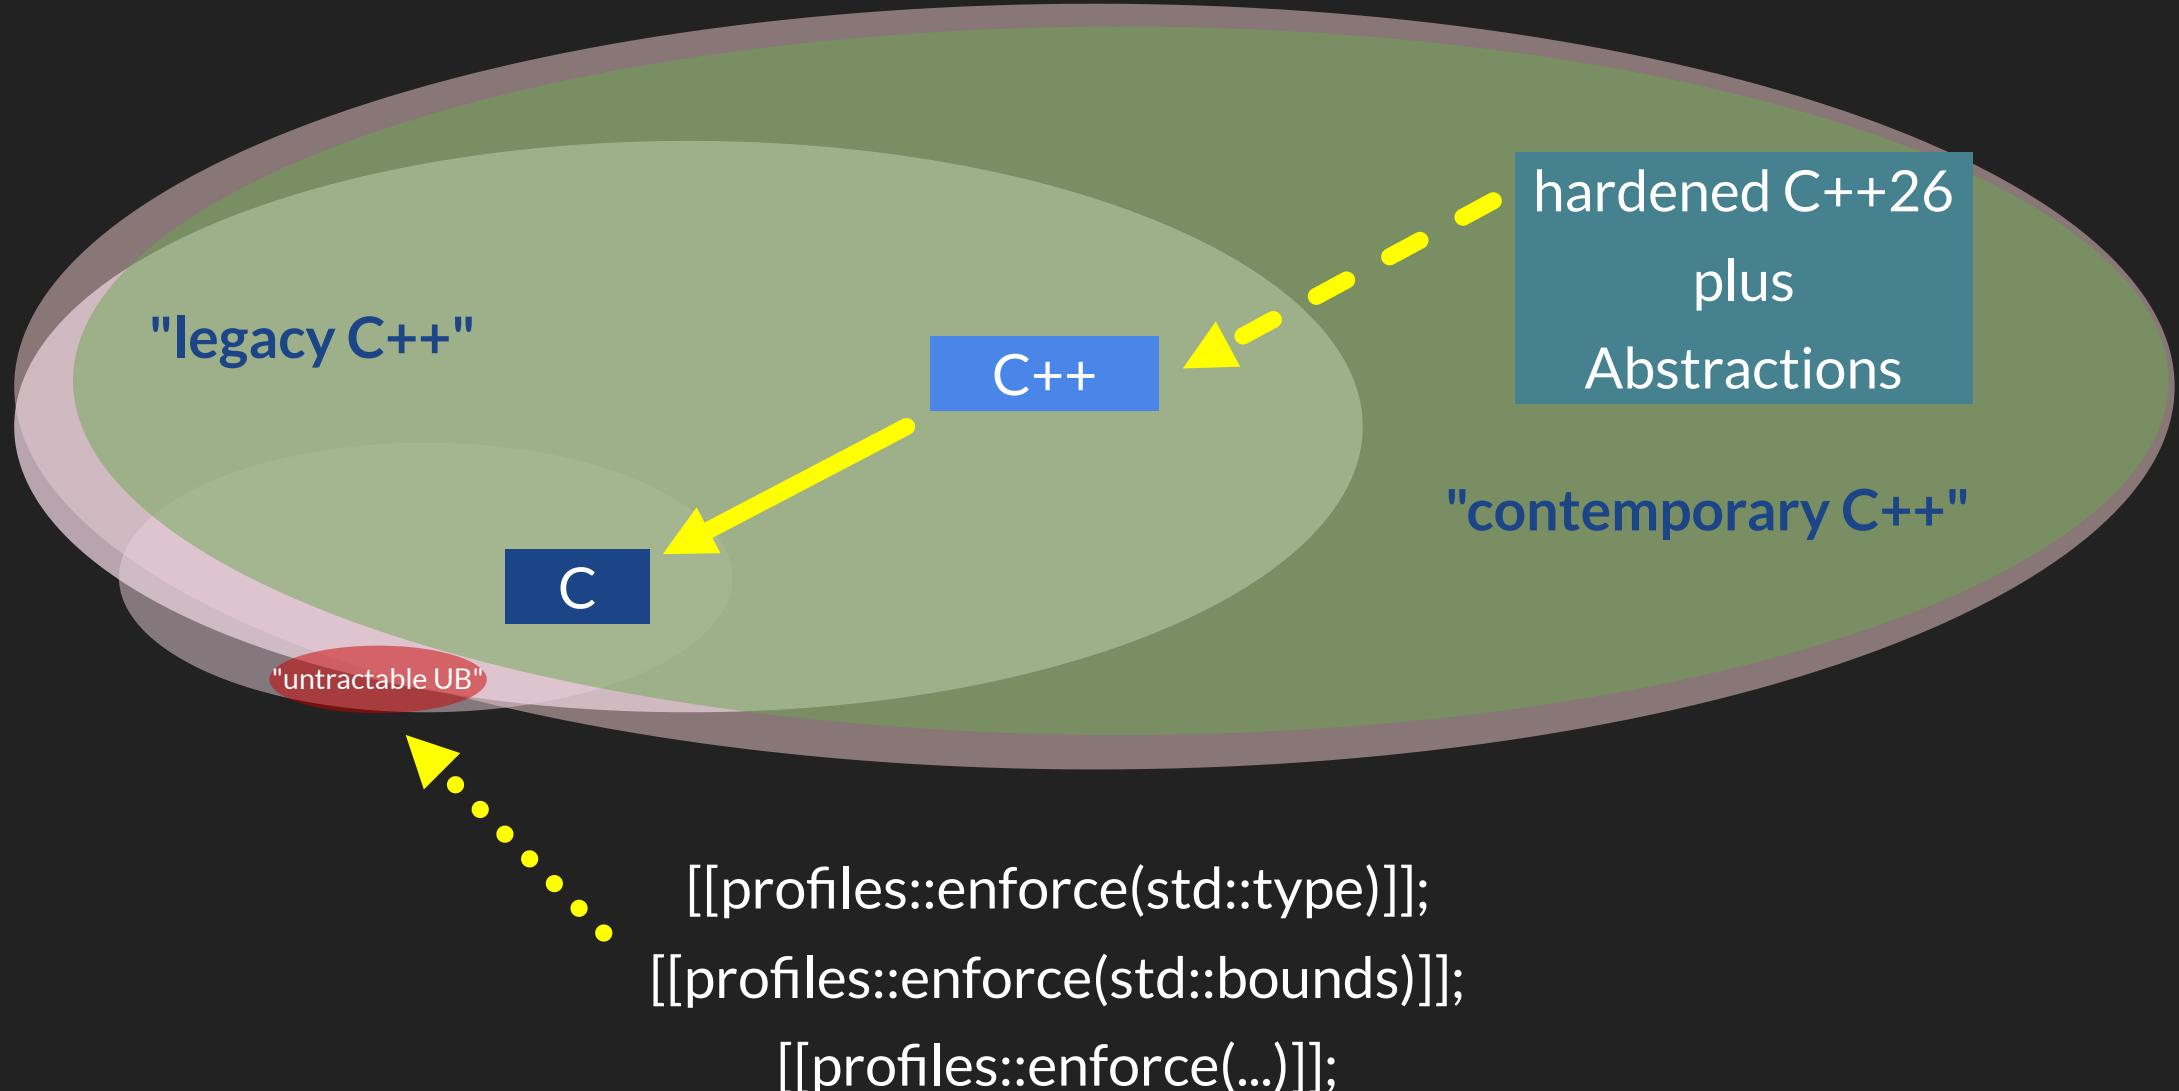
THE IDEA



Inherited UB taints the entire language

C++ adds its own UB

THE IDEA



SUMMARY

- why this perceived push towards security?
- security versus functional safety
- implications of UB at runtime and compile-time
- compile-time testing
- hardening the C++ standard library
- C++ contracts
- C++ profiles
- path into the future

"Simplicity is prerequisite
for reliability"

Edsger W. Dijkstra

RESOURCES

- Living, up-to-date C++ standard (currently at C++26 Committee Draft)
- The Case for Memory Safe Roadmaps
- P2900 Contracts for C++
- P3081 Core safety profiles for C++26
- P3589 C++ Profiles: The Framework

Contact

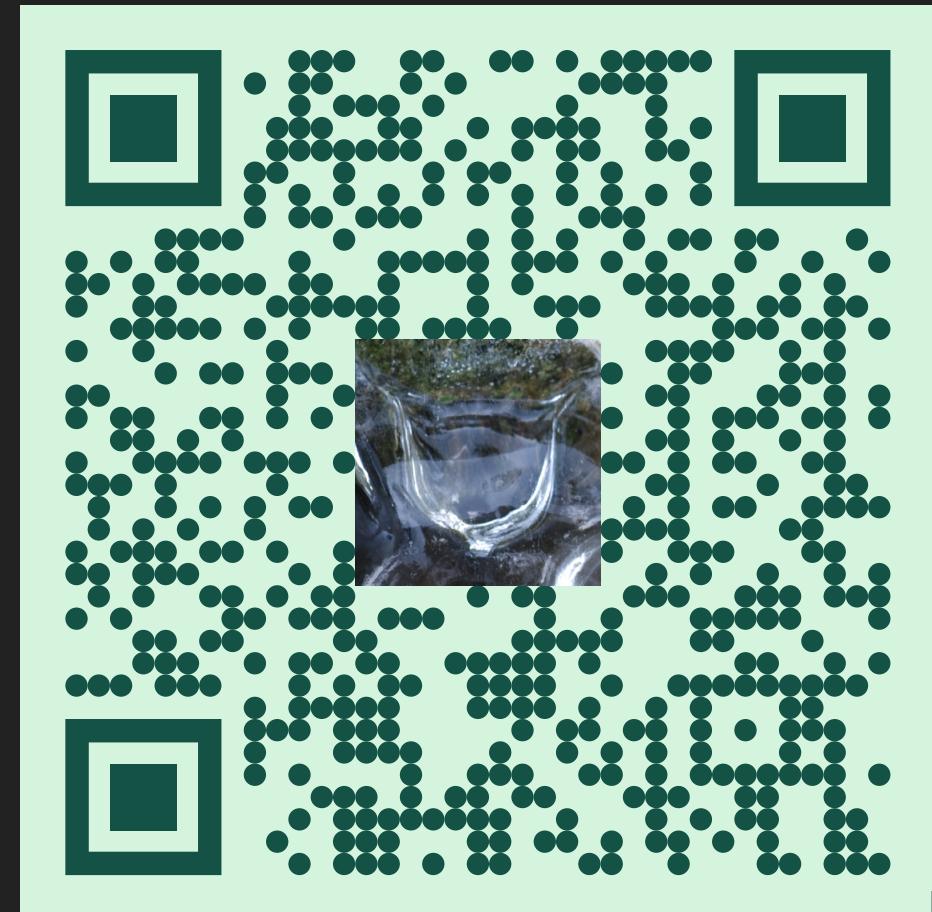
 dani@ngrt.de

 @DanielaKEngert@hachyderm.io

 DanielaE

 DanielaE

Images: courtesy of Matúš Chochlík





Ceterum censeo ABI esse frangendam